

Sri Sivasubramaniya Nadar College of Engineering, Chennai

(An Autonomous Institution affiliated to Anna University)

Machine Learning Laboratory

Experiment 9

Perceptron vs Multilayer Perceptron (A/B Experiment)

Name: MADHURI B

Register No.: 3122247001033

Aim and Objective

To compare the performance of a Single-Layer Perceptron Learning Algorithm (PLA) and a Multilayer Perceptron (MLP) for handwritten English character classification.

The objectives are:

- To implement a multiclass Perceptron Learning Algorithm.
- To implement an MLP with hidden layers and nonlinear activation functions.
- To tune the important parameters of the MLP.
- To compare different cost functions.
- To compare PLA and MLP using accuracy, precision, recall and F1-score.
- To analyse the models using confusion matrices and ROC curves.

Dataset Description

The dataset used for this experiment is the **Chars74K English Handwritten Characters Dataset (EnglishHnd.tgz)**.

The dataset contains handwritten English characters belonging to 62 classes. The classes include numbers, uppercase letters and lowercase letters.

- Number of classes: 62
- Total number of images: 3410

- Samples per class: 55
- Classes: 0–9, A–Z and a–z

Preprocessing Steps

The following preprocessing steps were used:

1. The handwritten character images were loaded from the dataset.
2. The images were converted into grayscale.
3. Each image was resized to 28×28 pixels.
4. Pixel values were normalized between 0 and 1.
5. The images were flattened into 784 features.
6. The dataset was divided into training, validation and testing sets.

The data obtained after splitting was:

- Training set: 2182 samples
- Validation set: 546 samples
- Test set: 682 samples

Perceptron Learning Algorithm (PLA)

A multiclass Perceptron was implemented from scratch. A separate weight vector was maintained for each of the 62 classes.

For each input image, the class having the highest score was taken as the predicted class. If the prediction was wrong, the weights of the actual and predicted classes were updated.

The learning rate used was 0.01 and the model was trained for 15 epochs.

The update was based on the Perceptron learning rule:

$$w_{t+1} = w_t + \eta(y - \hat{y})x$$

where η is the learning rate, y is the actual output and $\hat{y}$ is the predicted output.

PLA Results

Metric	PLA
Accuracy	0.1422
Precision	0.2574
Recall	0.1422
F1-score	0.1307

Table 1: PLA performance

The training error reduced from 0.9803 in the first epoch to 0.7832 in the fifteenth epoch.

Multilayer Perceptron (MLP)

The MLP was used with an input layer, hidden layer(s) and a 62-class output layer.

The output layer used softmax activation because the problem contains multiple character classes.

Different activation functions, optimizers, learning rates, hidden layer structures and batch sizes were tested.

Cost Function Comparison

Two cost functions were compared using the same basic MLP setup:

- Categorical Cross-Entropy
- Mean Squared Error (MSE)

Cost Function	Best Validation Accuracy
Categorical Cross-Entropy	0.1337
MSE	0.1062

Table 2: Cost function comparison

Categorical Cross-Entropy gave better validation accuracy in the comparison and was used for the final MLP.

Hyperparameter Tuning

Different combinations of hidden layers, activation functions, optimizers, learning rates and batch sizes were tested.

Hidden Layers	Activation	Optimizer	LR	Batch Size	Val. Accuracy
(128,)	Tanh	Adam	0.001	32	0.1172
(128,64)	Tanh	Adam	0.001	32	0.1136
(128,)	ReLU	Adam	0.0005	64	0.0623
(128,64)	ReLU	Adam	0.001	32	0.0568
(64,)	ReLU	SGD	0.01	32	0.0330
(128,)	ReLU	Adam	0.001	32	0.0165
(128,)	Tanh	GD	0.001	2182	0.0092

Table 3: Hyperparameter tuning results

The selected configuration was:

- Hidden layers: (128,)
- Activation function: Tanh
- Optimizer: Adam
- Learning rate: 0.001
- Batch size: 32

The best validation accuracy obtained during tuning was 0.1172.

Final MLP Results

The selected MLP configuration was trained for 25 epochs.

Metric	Tuned MLP
Accuracy	0.2331
Precision	0.2423
Recall	0.2331
F1-score	0.2074

Table 4: Final MLP performance

The training loss reduced from 4.2662 in the first epoch to 2.9228 in the final epoch.

A/B Comparison

The final performance of PLA and the tuned MLP is given below.

Model	Accuracy	Precision	Recall	F1-score
PLA	0.1422	0.2574	0.1422	0.1307
Tuned MLP	0.2331	0.2423	0.2331	0.2074

Table 5: PLA and MLP comparison

The MLP obtained higher accuracy and F1-score than the PLA on the test data.

Confusion Matrix and ROC

Confusion matrices were generated for both PLA and MLP to study the class-wise predictions.

ROC curves were also generated using micro-average and macro-average methods.

Model	Micro AUC	Macro AUC
PLA	0.7797	0.8543
MLP	0.8577	0.8609

Table 6: ROC AUC values

Experimental Results

The final test results obtained from the experiment are:

Model	Accuracy	Precision	Recall	F1-score
PLA	0.1422	0.2574	0.1422	0.1307
Tuned MLP	0.2331	0.2423	0.2331	0.2074

Table 7: Final experimental results

Experiment 9: Perceptron vs Multilayer Perceptron (A/B Experiment)

Dataset: Chars74K English Handwritten Characters (`EnglishHnd.tgz`)

Objective: Compare a Single-Layer Perceptron (PLA) with a tuned Multilayer Perceptron (MLP) for handwritten English character classification.

1. Import Libraries

```
In [1]: import os
import tarfile
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

from PIL import Image
from pathlib import Path

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, label_binarize
from sklearn.metrics import accuracy_score, precision_score, recall_score
from sklearn.metrics import confusion_matrix, classification_report, roc_

import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense, Input
from tensorflow.keras.optimizers import SGD, Adam

np.random.seed(42)
tf.random.set_seed(42)

print("TensorFlow version:", tf.__version__)
```

TensorFlow version: 2.20.0

2. Upload and Extract the Dataset

```
In [2]: try:
    from google.colab import files

    uploaded = files.upload()
    uploaded_name = next(iter(uploaded))
```

```

DATA_ROOT = Path("/content/chars74k_data")
DATA_ROOT.mkdir(exist_ok=True)

archive_path = Path("/content") / uploaded_name

if uploaded_name.endswith((".tgz", ".tar.gz", ".tar")):
    with tarfile.open(archive_path, "r:") as tar:
        tar.extractall(DATA_ROOT)
else:
    raise ValueError("Please upload EnglishHnd.tgz")

except ImportError:
    # For local Jupyter, put EnglishHnd.tgz in the notebook folder.
    DATA_ROOT = Path(".")
    archive_path = Path("EnglishHnd.tgz")

    if archive_path.exists():
        with tarfile.open(archive_path, "r:") as tar:
            tar.extractall(DATA_ROOT)

print("Dataset extraction completed.")

```

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving EnglishHnd.tgz to EnglishHnd.tgz

```

/tmp/ipykernel_507/2550353038.py:14: DeprecationWarning: Python 3.14 will, by default, filter extracted tar archives and reject files or modify their metadata. Use the filter argument to control this behavior.
    tar.extractall(DATA_ROOT)
Dataset extraction completed.

```

```

In [3]: img_root = None

for p in DATA_ROOT.rglob("Img"):
    if p.is_dir():
        img_root = p
        break

if img_root is None:
    raise FileNotFoundError("Could not find the Img folder inside EnglishHnd")

sample_dirs = sorted([p for p in img_root.iterdir() if p.is_dir() and p.name != "."])

print("Image folder:", img_root)
print("Number of classes:", len(sample_dirs))
print("First few classes:", [p.name for p in sample_dirs[:5]])
print("Last few classes:", [p.name for p in sample_dirs[-5:]])

```

```

Image folder: /content/chars74k_data/English/Hnd/Img
Number of classes: 62
First few classes: ['Sample001', 'Sample002', 'Sample003', 'Sample004', 'Sample005']
Last few classes: ['Sample058', 'Sample059', 'Sample060', 'Sample061', 'Sample062']

```

3. Dataset Description

```
In [4]: class_names = list("0123456789ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz")

print("Expected classes:", len(class_names))
print("Expected images:", 62 * 55)

png_files = list(img_root.rglob("*.png"))
print("PNG images found:", len(png_files))
```

Expected classes: 62
Expected images: 3410
PNG images found: 3410

4. Display Some Sample Images

```
In [5]: fig, axes = plt.subplots(2, 5, figsize=(10, 4))

for ax, path in zip(axes.ravel(), png_files[:10]):
    img = Image.open(path).convert("L")
    ax.imshow(img, cmap="gray")
    ax.set_title(path.parent.name)
    ax.axis("off")

plt.tight_layout()
plt.show()
```

Sample021



Sample021



Sample021



Sample021



Sample021



Sample021



Sample021



Sample021



Sample021



Sample021



5. Preprocessing

```
In [6]: IMG_SIZE = 28

X = []
y = []
```

```

for class_index, class_dir in enumerate(sample_dirs):
    image_files = sorted(class_dir.glob("*.png"))

    for image_path in image_files:
        img = Image.open(image_path).convert("L")
        img = img.resize((IMG_SIZE, IMG_SIZE))
        arr = np.asarray(img, dtype=np.float32) / 255.0

        X.append(arr.flatten())
        y.append(class_index)

X = np.array(X, dtype=np.float32)
y = np.array(y, dtype=np.int64)

print("X shape:", X.shape)
print("y shape:", y.shape)
print("Number of classes:", len(np.unique(y)))
print("Minimum pixel value:", X.min())
print("Maximum pixel value:", X.max())

```

```

X shape: (3410, 784)
y shape: (3410,)
Number of classes: 62
Minimum pixel value: 0.0
Maximum pixel value: 1.0

```

```

In [7]: label_names = class_names[:len(np.unique(y))]

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.20,
    random_state=42,
    stratify=y
)

X_train, X_val, y_train, y_val = train_test_split(
    X_train, y_train,
    test_size=0.20,
    random_state=42,
    stratify=y_train
)

print("Training:", X_train.shape)
print("Validation:", X_val.shape)
print("Testing:", X_test.shape)

```

```

Training: (2182, 784)
Validation: (546, 784)
Testing: (682, 784)

```

6. Model A — Multiclass Perceptron Learning Algorithm (PLA)

```
In [8]: class MulticlassPerceptron:
    def __init__(self, n_classes, n_features, learning_rate=0.01, epochs=
        self.n_classes = n_classes
        self.n_features = n_features
        self.learning_rate = learning_rate
        self.epochs = epochs
        self.W = np.zeros((n_classes, n_features), dtype=np.float32)
        self.b = np.zeros(n_classes, dtype=np.float32)
        self.error_history = []

    def predict(self, X):
        scores = X @ self.W.T + self.b
        return np.argmax(scores, axis=1)

    def decision_function(self, X):
        return X @ self.W.T + self.b

    def fit(self, X, y):
        for epoch in range(self.epochs):
            errors = 0
            indices = np.random.permutation(len(X))

            for i in indices:
                scores = X[i] @ self.W.T + self.b
                pred = np.argmax(scores)

                if pred != y[i]:
                    self.W[y[i]] += self.learning_rate * X[i]
                    self.b[y[i]] += self.learning_rate

                    self.W[pred] -= self.learning_rate * X[i]
                    self.b[pred] -= self.learning_rate

                errors += 1

            error_rate = errors / len(X)
            self.error_history.append(error_rate)

            print(f"Epoch {epoch+1:2d}/{self.epochs} - error rate: {error

        return self
```

```
In [9]: pla = MulticlassPerceptron(
    n_classes=len(label_names),
    n_features=X_train.shape[1],
    learning_rate=0.01,
    epochs=15
)

pla.fit(X_train, y_train)

pla_pred = pla.predict(X_test)
```

```

pla_accuracy = accuracy_score(y_test, pla_pred)
pla_precision = precision_score(y_test, pla_pred, average="macro", zero_d
pla_recall = recall_score(y_test, pla_pred, average="macro", zero_divisio
pla_f1 = f1_score(y_test, pla_pred, average="macro", zero_division=0)

print("\nPLA Results")
print("Accuracy :", round(pla_accuracy, 4))
print("Precision:", round(pla_precision, 4))
print("Recall   :", round(pla_recall, 4))
print("F1-score :", round(pla_f1, 4))

```

```

Epoch 1/15 - error rate: 0.9803
Epoch 2/15 - error rate: 0.9583
Epoch 3/15 - error rate: 0.9248
Epoch 4/15 - error rate: 0.9047
Epoch 5/15 - error rate: 0.8932
Epoch 6/15 - error rate: 0.8845
Epoch 7/15 - error rate: 0.8639
Epoch 8/15 - error rate: 0.8515
Epoch 9/15 - error rate: 0.8318
Epoch 10/15 - error rate: 0.8185
Epoch 11/15 - error rate: 0.8135
Epoch 12/15 - error rate: 0.8194
Epoch 13/15 - error rate: 0.8176
Epoch 14/15 - error rate: 0.8126
Epoch 15/15 - error rate: 0.7832

```

```


PLA Results
Accuracy : 0.1422
Precision: 0.2574
Recall   : 0.1422
F1-score : 0.1307

```

```

In [10... plt.figure(figsize=(7, 4))
plt.plot(range(1, len(pla.error_history) + 1), pla.error_history, marker=
plt.xlabel("Epoch")
plt.ylabel("Training Error Rate")
plt.title("PLA Training Error vs Epochs")
plt.grid(True)
plt.show()

```



7. Model B — Multilayer Perceptron (MLP)

```
In [33... def build_mlp(hidden_layers=(128,), activation="relu", optimizer_name="ad
            learning_rate=0.001, loss_function="sparse_categorical_cros

    model = Sequential()
    model.add(Input(shape=(X_train.shape[1],)))

    for units in hidden_layers:
        model.add(Dense(units, activation=activation))

    model.add(Dense(len(label_names), activation="softmax"))

    if optimizer_name.lower() in ["gd", "sgd"]:
        optimizer = SGD(learning_rate=learning_rate)
    else:
        optimizer = Adam(learning_rate=learning_rate)

    model.compile(
        optimizer=optimizer,
        loss=loss_function,
        metrics=["accuracy"]
    )

    return model
```

```
In [34... # Cost function comparison
```

```

from tensorflow.keras.utils import to_categorical

y_train_cat = to_categorical(y_train, num_classes=len(label_names))
y_val_cat = to_categorical(y_val, num_classes=len(label_names))

loss_results = []

for loss_name in ["categorical_crossentropy", "mse"]:

    model = Sequential([
        Input(shape=(X_train.shape[1],)),
        Dense(128, activation="tanh"),
        Dense(len(label_names), activation="softmax")
    ])

    model.compile(
        optimizer=Adam(learning_rate=0.001),
        loss=loss_name,
        metrics=["accuracy"]
    )

    history = model.fit(
        X_train, y_train_cat,
        validation_data=(X_val, y_val_cat),
        epochs=10,
        batch_size=32,
        verbose=0
    )

    loss_results.append({
        "Cost Function": loss_name,
        "Best Validation Accuracy": max(history.history["val_accuracy"])
    })

loss_df = pd.DataFrame(loss_results)

print("Cost Function Comparison")
display(loss_df)

```

Cost Function Comparison

	Cost Function	Best Validation Accuracy
0	categorical_crossentropy	0.133700
1	mse	0.106227

8. Hyperparameter Tuning

In [35... configs = [

```

    {"hidden_layers": (64,), "activation": "relu", "optimizer": "sgd", "l
    {"hidden_layers": (128,), "activation": "relu", "optimizer": "adam",

```

```

{"hidden_layers": (128,), "activation": "tanh", "optimizer": "adam",
{"hidden_layers": (128,), "activation": "relu", "optimizer": "adam",
{"hidden_layers": (128, 64), "activation": "relu", "optimizer": "adam",
{"hidden_layers": (128, 64), "activation": "tanh", "optimizer": "adam",
    {"hidden_layers": (128,), "activation": "tanh", "optimizer": "sgd",
]

tuning_results = []

for i, cfg in enumerate(configs, start=1):
    print(f"\nConfiguration {i}/{len(configs)}")
    print(cfg)

    model = build_mlp(
        hidden_layers=cfg["hidden_layers"],
        activation=cfg["activation"],
        optimizer_name=cfg["optimizer"],
        learning_rate=cfg["learning_rate"]
    )

    history = model.fit(
        X_train, y_train,
        validation_data=(X_val, y_val),
        epochs=12,
        batch_size=cfg["batch_size"],
        verbose=0
    )

    best_val_acc = max(history.history["val_accuracy"])
    final_val_acc = history.history["val_accuracy"][-1]

    tuning_results.append({
        **cfg,
        "best_validation_accuracy": best_val_acc,
        "final_validation_accuracy": final_val_acc
    })

    print("Best validation accuracy:", round(best_val_acc, 4))

tuning_df = pd.DataFrame(tuning_results)
tuning_df = tuning_df.sort_values("best_validation_accuracy", ascending=False)
tuning_df

```

Configuration 1/7
{'hidden_layers': (64,), 'activation': 'relu', 'optimizer': 'sgd', 'learning_rate': 0.01, 'batch_size': 32}
Best validation accuracy: 0.033

Configuration 2/7
{'hidden_layers': (128,), 'activation': 'relu', 'optimizer': 'adam', 'learning_rate': 0.001, 'batch_size': 32}
Best validation accuracy: 0.0165

Configuration 3/7
{'hidden_layers': (128,), 'activation': 'tanh', 'optimizer': 'adam', 'learning_rate': 0.001, 'batch_size': 32}
Best validation accuracy: 0.1172

Configuration 4/7
{'hidden_layers': (128,), 'activation': 'relu', 'optimizer': 'adam', 'learning_rate': 0.0005, 'batch_size': 64}
Best validation accuracy: 0.0623

Configuration 5/7
{'hidden_layers': (128, 64), 'activation': 'relu', 'optimizer': 'adam', 'learning_rate': 0.001, 'batch_size': 32}
Best validation accuracy: 0.0568

Configuration 6/7
{'hidden_layers': (128, 64), 'activation': 'tanh', 'optimizer': 'adam', 'learning_rate': 0.001, 'batch_size': 32}
Best validation accuracy: 0.1136

Configuration 7/7
{'hidden_layers': (128,), 'activation': 'tanh', 'optimizer': 'sgd', 'learning_rate': 0.001, 'batch_size': 2182}
Best validation accuracy: 0.0092

```
Out[35]:
```

	hidden_layers	activation	optimizer	learning_rate	batch_size	best_validation_accuracy
0	(128,)	tanh	adam	0.0010	32	0.117216
1	(128, 64)	tanh	adam	0.0010	32	0.113553
2	(128,)	relu	adam	0.0005	64	0.062271
3	(128, 64)	relu	adam	0.0010	32	0.056777
4	(64,)	relu	sgd	0.0100	32	0.032967
5	(128,)	relu	adam	0.0010	32	0.016484
6	(128,)	tanh	sgd	0.0010	2182	0.009158

9. Select the Best MLP Configuration

```
In [36]: best_cfg = tuning_df.iloc[0].to_dict()
```

```

print("Selected configuration:")
print("Hidden layers :", best_cfg["hidden_layers"])
print("Activation      :", best_cfg["activation"])
print("Optimizer       :", best_cfg["optimizer"])
print("Learning rate  :", best_cfg["learning_rate"])
print("Batch size      :", int(best_cfg["batch_size"]))
print("Best val acc   :", round(best_cfg["best_validation_accuracy"], 4))

```

```

Selected configuration:
Hidden layers : (128,)
Activation      : tanh
Optimizer       : adam
Learning rate  : 0.001
Batch size     : 32
Best val acc   : 0.1172

```

10. Train the Final Tuned MLP

```

In [37... final_mlp = build_mlp(
    hidden_layers=tuple(best_cfg["hidden_layers"]),
    activation=best_cfg["activation"],
    optimizer_name=best_cfg["optimizer"],
    learning_rate=float(best_cfg["learning_rate"])
)

final_history = final_mlp.fit(
    X_train, y_train,
    validation_data=(X_val, y_val),
    epochs=25,
    batch_size=int(best_cfg["batch_size"]),
    verbose=1
)

```

Epoch 1/25
69/69  **2s** 11ms/step - accuracy: 0.0165 - loss: 4.2662
- val_accuracy: 0.0183 - val_loss: 4.1372
Epoch 2/25
69/69  **0s** 5ms/step - accuracy: 0.0151 - loss: 4.1664
- val_accuracy: 0.0256 - val_loss: 4.1020
Epoch 3/25
69/69  **0s** 4ms/step - accuracy: 0.0247 - loss: 4.1258
- val_accuracy: 0.0458 - val_loss: 4.0706
Epoch 4/25
69/69  **0s** 4ms/step - accuracy: 0.0325 - loss: 4.0658
- val_accuracy: 0.0495 - val_loss: 4.0078
Epoch 5/25
69/69  **0s** 4ms/step - accuracy: 0.0518 - loss: 3.9878
- val_accuracy: 0.0604 - val_loss: 3.9264
Epoch 6/25
69/69  **0s** 4ms/step - accuracy: 0.0701 - loss: 3.9043
- val_accuracy: 0.0696 - val_loss: 3.8496
Epoch 7/25
69/69  **0s** 4ms/step - accuracy: 0.0944 - loss: 3.8183
- val_accuracy: 0.1026 - val_loss: 3.7653
Epoch 8/25
69/69  **0s** 4ms/step - accuracy: 0.1091 - loss: 3.7310
- val_accuracy: 0.1245 - val_loss: 3.6907
Epoch 9/25
69/69  **0s** 4ms/step - accuracy: 0.1306 - loss: 3.6475
- val_accuracy: 0.1136 - val_loss: 3.6297
Epoch 10/25
69/69  **0s** 4ms/step - accuracy: 0.1370 - loss: 3.5698
- val_accuracy: 0.1502 - val_loss: 3.5594
Epoch 11/25
69/69  **0s** 4ms/step - accuracy: 0.1577 - loss: 3.5002
- val_accuracy: 0.1703 - val_loss: 3.5012
Epoch 12/25
69/69  **0s** 4ms/step - accuracy: 0.1641 - loss: 3.4343
- val_accuracy: 0.1886 - val_loss: 3.4532
Epoch 13/25
69/69  **0s** 4ms/step - accuracy: 0.1783 - loss: 3.3756
- val_accuracy: 0.1941 - val_loss: 3.4121
Epoch 14/25
69/69  **0s** 4ms/step - accuracy: 0.1966 - loss: 3.3210
- val_accuracy: 0.1996 - val_loss: 3.3722
Epoch 15/25
69/69  **0s** 5ms/step - accuracy: 0.2016 - loss: 3.2699
- val_accuracy: 0.2033 - val_loss: 3.3392
Epoch 16/25
69/69  **0s** 4ms/step - accuracy: 0.2236 - loss: 3.2216
- val_accuracy: 0.2125 - val_loss: 3.3117
Epoch 17/25
69/69  **0s** 4ms/step - accuracy: 0.2241 - loss: 3.1794
- val_accuracy: 0.2198 - val_loss: 3.2908
Epoch 18/25
69/69  **0s** 4ms/step - accuracy: 0.2324 - loss: 3.1381
- val_accuracy: 0.2106 - val_loss: 3.2696
Epoch 19/25

```

69/69 ————— 0s 4ms/step - accuracy: 0.2397 - loss: 3.0997
- val_accuracy: 0.1941 - val_loss: 3.2606
Epoch 20/25
69/69 ————— 0s 4ms/step - accuracy: 0.2489 - loss: 3.0637
- val_accuracy: 0.2271 - val_loss: 3.2313
Epoch 21/25
69/69 ————— 0s 4ms/step - accuracy: 0.2608 - loss: 3.0311
- val_accuracy: 0.2271 - val_loss: 3.2130
Epoch 22/25
69/69 ————— 0s 4ms/step - accuracy: 0.2649 - loss: 3.0008
- val_accuracy: 0.2363 - val_loss: 3.1935
Epoch 23/25
69/69 ————— 0s 5ms/step - accuracy: 0.2676 - loss: 2.9709
- val_accuracy: 0.2399 - val_loss: 3.1751
Epoch 24/25
69/69 ————— 0s 4ms/step - accuracy: 0.2681 - loss: 2.9488
- val_accuracy: 0.2363 - val_loss: 3.1607
Epoch 25/25
69/69 ————— 0s 4ms/step - accuracy: 0.2722 - loss: 2.9228
- val_accuracy: 0.2527 - val_loss: 3.1304

```

```

In [38... mlp_prob = final_mlp.predict(X_test, verbose=0)
mlp_pred = np.argmax(mlp_prob, axis=1)

mlp_accuracy = accuracy_score(y_test, mlp_pred)
mlp_precision = precision_score(y_test, mlp_pred, average="macro", zero_d
mlp_recall = recall_score(y_test, mlp_pred, average="macro", zero_divisio
mlp_f1 = f1_score(y_test, mlp_pred, average="macro", zero_division=0)

print("\nTuned MLP Results")
print("Accuracy :", round(mlp_accuracy, 4))
print("Precision:", round(mlp_precision, 4))
print("Recall    :", round(mlp_recall, 4))
print("F1-score  :", round(mlp_f1, 4))


Tuned MLP Results
Accuracy : 0.2331
Precision: 0.2423
Recall    : 0.2331
F1-score  : 0.2074

```

```

In [39... plt.figure(figsize=(7, 4))
plt.plot(final_history.history["loss"], label="Training Loss")
plt.plot(final_history.history["val_loss"], label="Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.title("MLP Training Error vs Epochs")
plt.legend()
plt.grid(True)
plt.show()

```





11. Confusion Matrices

```
In [40... cm_pla = confusion_matrix(y_test, pla_pred)
cm_mlp = confusion_matrix(y_test, mlp_pred)

fig, ax = plt.subplots(figsize=(14, 12))
im = ax.imshow(cm_pla, interpolation="nearest", cmap="Blues")
ax.set_title("PLA Confusion Matrix")
ax.set_xlabel("Predicted Class")
ax.set_ylabel("True Class")
plt.colorbar(im, ax=ax)
plt.tight_layout()
plt.show()

fig, ax = plt.subplots(figsize=(14, 12))
im = ax.imshow(cm_mlp, interpolation="nearest", cmap="Blues")
ax.set_title("MLP Confusion Matrix")
ax.set_xlabel("Predicted Class")
ax.set_ylabel("True Class")
plt.colorbar(im, ax=ax)
plt.tight_layout()
plt.show()
```





12. ROC Curves — Micro and Macro Average

```
In [41... y_test_bin = label_binarize(y_test, classes=np.arange(len(label_names)))

pla_scores = pla.decision_function(X_test)

# Convert PLA scores to a probability-like scale for ROC calculation.
pla_exp = np.exp(pla_scores - np.max(pla_scores, axis=1, keepdims=True))
pla_prob = pla_exp / pla_exp.sum(axis=1, keepdims=True)

def plot_micro_macro_roc(y_true_bin, scores, title):
    fpr = {}
    tpr = {}
    roc_auc = {}
```

```

n_classes = y_true_bin.shape[1]

for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_true_bin[:, i], scores[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

fpr["micro"], tpr["micro"], _ = roc_curve(y_true_bin.ravel(), scores)
roc_auc["micro"] = auc(fpr["micro"], tpr["micro"])

all_fpr = np.unique(np.concatenate([fpr[i] for i in range(n_classes)]))
mean_tpr = np.zeros_like(all_fpr)

for i in range(n_classes):
    mean_tpr += np.interp(all_fpr, fpr[i], tpr[i])

mean_tpr /= n_classes
fpr["macro"] = all_fpr
tpr["macro"] = mean_tpr
roc_auc["macro"] = auc(fpr["macro"], tpr["macro"])

plt.figure(figsize=(7, 5))
plt.plot(fpr["micro"], tpr["micro"],
         label=f"Micro-average (AUC = {roc_auc['micro']:.3f})")
plt.plot(fpr["macro"], tpr["macro"],
         label=f"Macro-average (AUC = {roc_auc['macro']:.3f})")
plt.plot([0, 1], [0, 1], linestyle="--")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title(title)
plt.legend()
plt.grid(True)
plt.show()

return roc_auc


pla_roc_auc = plot_micro_macro_roc(
    y_test_bin, pla_prob, "PLA ROC Curve"
)


mlp_roc_auc = plot_micro_macro_roc(
    y_test_bin, mlp_prob, "MLP ROC Curve"
)

print("PLA Micro AUC:", round(pla_roc_auc["micro"], 4))
print("PLA Macro AUC:", round(pla_roc_auc["macro"], 4))
print("MLP Micro AUC:", round(mlp_roc_auc["micro"], 4))
print("MLP Macro AUC:", round(mlp_roc_auc["macro"], 4))

```

PLA ROC Curve





PLA Micro AUC: 0.7797
 PLA Macro AUC: 0.8543
 MLP Micro AUC: 0.8577
 MLP Macro AUC: 0.8609

13. Final A/B Comparison

```
In [42...] comparison = pd.DataFrame({
    "Model": ["PLA", "Tuned MLP"],
    "Accuracy": [pla_accuracy, mlp_accuracy],
    "Precision": [pla_precision, mlp_precision],
    "Recall": [pla_recall, mlp_recall],
    "F1-score": [pla_f1, mlp_f1]
})

comparison.round(4)
```

```
Out[42]:
```

	Model	Accuracy	Precision	Recall	F1-score
0	PLA	0.1422	0.2574	0.1422	0.1307
1	Tuned MLP	0.2331	0.2423	0.2331	0.2074

14. Observation

```
In [43... print("PLA Accuracy :", round(pla_accuracy, 4))
print("MLP Accuracy :", round(mlp_accuracy, 4))
print("PLA F1-score :", round(pla_f1, 4))
print("MLP F1-score :", round(mlp_f1, 4))
print()
print("Best MLP configuration:")
print(best_cfg)

PLA Accuracy : 0.1422
MLP Accuracy : 0.2331
PLA F1-score : 0.1307
MLP F1-score : 0.2074

Best MLP configuration:
{'hidden_layers': (128,), 'activation': 'tanh', 'optimizer': 'adam', 'learning_rate': 0.001, 'batch_size': 32, 'best_validation_accuracy': 0.11721611768007278, 'final_validation_accuracy': 0.11721611768007278}
```

In [43...

Learning Outcomes

- Understood the working of the Perceptron Learning Algorithm.
- Implemented a multiclass Perceptron from scratch.
- Understood the use of hidden layers and activation functions in MLP.
- Studied the effect of learning rate, batch size and optimizer.
- Compared different cost functions.
- Evaluated the models using accuracy, precision, recall and F1-score.
- Generated confusion matrices and ROC curves for model comparison.

Conclusion

The Perceptron and Multilayer Perceptron were implemented for handwritten English character classification using the Chars74K dataset.

The PLA obtained an accuracy of 0.1422, while the final tuned MLP obtained an accuracy of 0.2331. The F1-score of the MLP was 0.2074 compared with 0.1307 for the PLA.

The experiment helped in understanding the difference between a simple linear Perceptron and an MLP with hidden layers and nonlinear activation functions. Hyperparameter tuning was also used to select the MLP configuration used for the final test.